



INTRODUCTION TO AI – UNSUPERVISED LEARNING

Philip Mortimer



Supervised learning



Unsupervised learning

RECAP

- Reinforced knowledge of supervised learning by building an ANN to identify cancerous tumours
- Discussed a modern variant of neural networks, CNN's
- CNN's use feature extraction to learn high level details
- This allows AI systems to be scaled for large inputs (e.g. high quality images)
- We discussed TensorFlow, a popular machine learning library
- Used TensorFlow to implement CNN and ANN to tackle handwritten digit recognition. We used MNIST dataset

UNSUPERVISED LEARNING

- Supervised learning involves labelled datasets. These labels often have to be produced by people
- As machine learning requires large amounts of data to be effective, it is often very expensive to label the data required
- This makes supervised learning challenging and costly to implement
- It also means that the strength of a model is limited by the accuracy of the labels
- This makes AI prone to human error and bias

UNSUPERVISED LEARNING

- Therefore, it is often much more appropriate to use unlabelled data
- This is often more cost effective
- There are also a wide range of systems where labelling data is simply inappropriate
- For example, we may have a social media website with a large number of users
- We may want to make an AI that recommends content that they are likely to find interesting
- We as outside observers have no efficient way to label these users' interests (we don't know what that person actually likes)
- However, a computer might be able to look at content that similar people like, and deduce that the user is likely to enjoy similar content


UNSUPERVISED LEARNING

- With supervised learning, there are a number of different algorithms to use (with gradient descent simply being the most common one)
- Unsupervised learning is a bucket term that encapsulates a range of possible algorithms
- Arguably the most popular one is k-means clustering

K-MEANS CLUSTERING


- The idea behind K-means clustering is that we have a set of datapoints and we want to group these data points into k clusters.
- The idea is that each of these clusters should represent similar datapoints
- The programmer specifies the number of different clusters that the data should be broken into (i.e. the value of k)

Before K-Means



K-Means

After K-Means




ALGORITHM

- We select k random points to be the initial “centroids” of each cluster
- We assign all points to one of these clusters by choosing the closest centroid for each point
- Now we can recompute the centroid of the cluster by aggregating the distance of all points
- We repeat this process with the newly calculated centroid
- This process occurs until the centroids no longer change (or until a specified number of iterations has passed)

K-MEANS

- STEP I – Choose K random centroids
- In this case, $k = 2$




Credit:

<https://www.analyticsvidhya.com/blog/2019/08/comprehensive-guide-k-means-clustering/>

K-MEANS

- STEP 2 – Assign all points to one of the k clusters based on which centroid they are closest to




Credit:

<https://www.analyticsvidhya.com/blog/2019/08/comprehensive-guide-k-means-clustering/>

K-MEANS

- STEP 3 – Calculate new centroids by taking that the aggregate location of the points within each cluster




Credit:

<https://www.analyticsvidhya.com/blog/2019/08/comprehensive-guide-k-means-clustering/>

K-MEANS

- Repeat Steps 2 and 3 until :
 1. The recalculated centroids are the same as the current centroids
 2. A fixed number of iterations has passed(i.e., we've spent enough time computing the solution)
 3. All points remain in same cluster over multiple iterations
- This gives us our clusters



Credit:

<https://www.analyticsvidhya.com/blog/2019/08/comprehensive-guide-k-means-clustering/>

K-MEANS

- K-Means enables us to group unlabelled data into categories that share common traits
- K-means essentially draws k hyperspheres which have the same dimensionality as the number of datapoint attributes used
- This has a very large number of practical applications

K-MEANS

- Imagine we are a bank with 5 million customers
- Let's say we wish to give each customer an offer to acquire a credit card
- Clearly, different customers will want different things from a credit card
- We can't sit down and decide what offer to make to each customer individually
- Hence, we could make 3 different offer packages targeted at different customers
- We may want one package to target people based on income. So we could split our packages into three classes: "high income, medium income, low income"
- Using records of our customers, we may decide to cluster them based on the following attributes: "amount in bank, value of outstanding loans and bank balance five years ago"
- This gives us a three dimensional space of five million datapoints

K-MEANS


- We then run the k means algorithm using $k = 3$ to cluster our data
- We know have 5 million customers neatly divided and can target out offers accordingly
- Note of course, that k means tells us nothing about **what** each cluster actually means. We need some way of figuring this out ourselves
- This could be achieved by human verification, our smart algorithms


K-MEANS

- In many cases, we won't know what value of k is most suitable
- A method known as the “elbow method” is often used. It involves choosing the value of k at which the distance between points in the cluster stops decreasing quickly

ELBOW METHOD

- We can see that at $k = 4$, the average distance within a cluster stops decreasing as rapidly
- At this point, the points have been grouped effectively
- We don't want k to be too large, as more groups often means more unneeded complexity





K-MEANS




- Many algorithms have already been written and tested by other programmers. We can simply reuse these if we wish (which saves us developing from scratch).

```
from sklearn.cluster import KMeans
```

```
def kMeansModel(X):  
    #X is data  
    model = KMeans(n_clusters=5)  
    model.fit(X)  
    return model
```

K-MEANS

- I've coded my own k-means learning model that clusters data into groups
- Let's look at the code. The code can be found here:
<https://github.com/philipmortimer/AI-Course/blob/main/Programs/Part%206/kMeans/main.py>



AUTO-ENCODERS

- Unsupervised learning covers a wide range of techniques and algorithms
- Image compression often uses a neural network structure known as an autoencoder
- These networks compress images by reducing their dimensionality by extracting the key features
- Auto-encoders are often trained on specific image types (e.g. only images of cats) in order to make them more effective at dimensionality reduction


AUTO-ENCODERS

- Auto-encoders require an unlabelled dataset (e.g. of images)
- We use supervised learning techniques to train the network (i.e. gradient descent)
- We have two networks, namely an “encoder” and a “decoder”
- The encoder network converts the image into an abstract representation that uses less data than the original image
- The decoder converts the compressed format into the original image
- By combining the two components, we have a network that can use data both as input and as a label
- Sometimes the input image is blurred slightly to add “noise” to the input

Autoencoders

Encoder

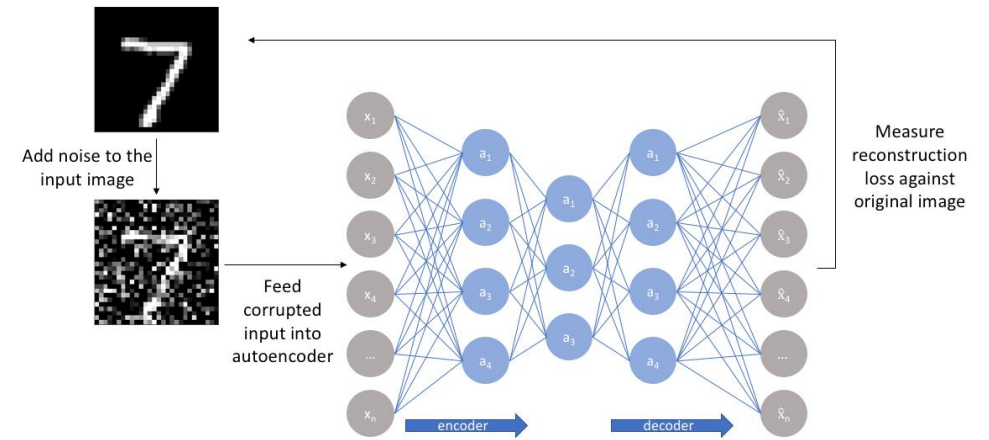
Decoder



Input

Latent Space
Representation

Output



AUTO-ENCODER

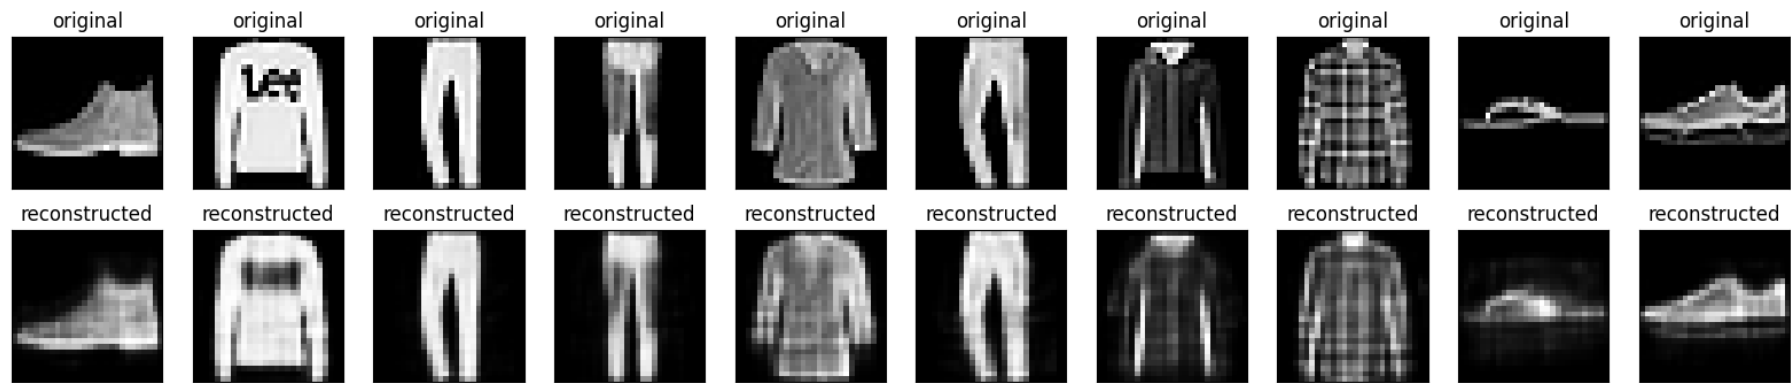
- As the compressed representation is dimensionally smaller than the input, the network must attempt to find a smaller representation of the image that allows the image to be best recreated
- Autoencoders are an example of lossy compression
- Autoencoders are often used to reduce the dimensionality of large datasets to make them easier to use (e.g. for machine learning techniques)
- For example, autoencoders have been applied to allow for more efficient representations of chess boards

AUTO-ENCODERS

- Training of auto-encoders is expensive
- The problem with auto-encoders is choosing the correct network layout (and other relevant supervised learning hyperparameters) to maximise compression, reconstruction accuracy and to minimise compute time
- This is often very difficult and time consuming

AUTO-ENCODERS


- I've trained my own autoencoder using TensorFlow
- This autoencoder takes in a 784 pixel image of a clothing item. This is then compressed to 64 pixels using the auto-encoder
- First, let's discuss the encoder. The encoder must take a 784 pixel input and convert it to a 64 value output. To achieve this, I have used no hidden layers (to maximise speed)
- The output layer contains 64 neurons.
- I have chosen to use the sigmoid activation
- For the decoder network, we must convert 64 values into 784 values
- I have once more elected to use no hidden layers to speed up computation
- Let's look at code! <https://github.com/philipmortimer/AI-Course/blob/main/Programs/Part%206/autoencoder/main.py>



AUTO-ENCODERS

AUTO-ENCODERS

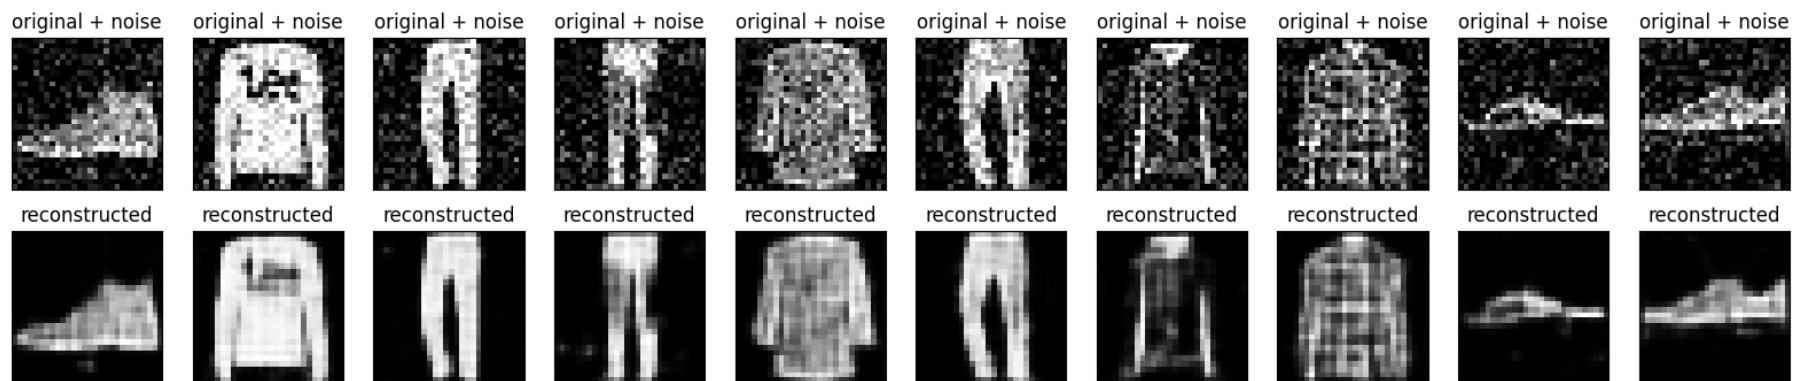
- Next, let's attempt to train an auto-encoder that removes noise from images
- To do this, we programmatically generate “noise” which has the effect of blurring our images somewhat
- We want to take this blurred image as input and return the unblurred image
- This is often used for post-processing of photos
- Additionally, it is used to increase the dataset size for traditional compression based autoencoders



NOISY IMAGES

DE-NOISING AUTOENCODER

- Notice, how we are mimicking supervised learning techniques.
- This system is unsupervised however, as it is using unlabelled data as both input and output
- For our de-noising autoencoder, I have decided to use convolutional layers in order to enable effective feature extraction
- Typically, convolutional approaches are more versatile and effective than just using ANN's
- Let's look at code! <https://github.com/philipmortimer/AI-Course/blob/main/Programs/Part%206/autoencoder/denoiseAutencoder.py>



DE-NOISED IMAGES

PRINCIPAL COMPONENT ANALYSIS (PCA)

- Principal component analysis (PCA) is another technique commonly used for dimension reduction
- Let's say we have a database containing millions of entries
- Each record stores three data attributes: x , y and z .
- Thus, each record is three dimensional
- PCA allows us to compress this dataset to 1d, 2d or 3d
- It does this by exploiting relationships between x , y and z .
- If it turns out that $x = 0.5y + 2z$, then we can reduce dimensionality as we clearly don't actually need to store x . It can simply be calculated from y and z .

PCA

- Before, doing anything else, it is very important that we standardise all of our attributes so that they have the same range
- This is important as PCA essentially captures maximum variance
- If one variable has range $[0, 255]$ whilst another has $[0, 1]$, the first variable will typically produce a much larger variance. This is misleading and hence both would need to be altered to fit same range (e.g. $[0, 1]$)

PCA

- Say we have records with n variables. This data is said to be n -dimensional
- We can then compute the covariance matrix for all possible combination of variable pairs.
- If the covariance is positive, it suggests that as one variable, increases, so does the other
- Similarly, if it's negative, it suggests that they are inversely correlated
- Small covariances suggest minimal correlation

$$\begin{bmatrix} Cov(x, x) & Cov(x, y) & Cov(x, z) \\ Cov(y, x) & Cov(y, y) & Cov(y, z) \\ Cov(z, x) & Cov(z, y) & Cov(z, z) \end{bmatrix}$$


Covariance matrix for 3d data

PCA

- The idea of PCA is to convert n variables into k variables where $k \leq n$
- To do this we will create k principal components. The idea is that we fit the most amount of information possible into the first principal component. Then as much into the second component and so on
- This allows us to use the first k principal components and ignore the rest, as these first k components will encode the most information

PCA

- This is an example for data which has 10 attributes
- One principal component attribute (which we must calculate) can explain about 40% of the data variance



Credit: <https://builtin.com/data-science/step-step-explanation-principal-component-analysis>

PCA

- If we calculate all of the eigenvectors of our covariance matrix (and normalise them where we needed), we get a set of vectors
- These vectors form an orthonormal basis
- Additionally, the corresponding eigenvalues indicate order of importance of each eigenvector (/principal component).
- This allows us to rank each principal component as seen on the previous page
- You don't have to understand the maths behind this, just that it allows us to effectively reduce dimensionality
- We can select the number of dimensions / accuracy we want
- It's important to note that the principal components lack interpretability

```
from sklearn.decomposition import PCA

def prin_comps(data, n_comps):
    pca = PCA(n_components=n_comps) #Reduces dimensionality to n_comps
    principalComponents = pca.fit_transform(data)
    return principalComponents
```

PCA

PCs # 0



PCs # 10



PCs # 20



PCs # 30



PCs # 40



PCs # 50



SUMMARY

- Unsupervised learning is a machine learning technique applied to datasets which are unlabelled
- There are a very wide range of algorithms used in unsupervised learning
- K-means clustering focuses on grouping datapoints into a certain number of groups. These clusters will group datapoints that are most similar to them
- Similarity between datapoints is calculated via the Euclidian distance in the vector space formed by the attributes of the points
- Autoencoders are often used to reduce dimensionality of data
- Autoencoders consist of an **encoder** neural network and a **decoder** neural network

SUMMARY

- The encoder network reduces the dimensionality of the network to a compressed format
- The decoder network takes this compressed format and attempts to convert it back to the original datapoint
- This can often be thought of as a feature extraction network
- Principal Component Analysis (PCA) is another dimensionality reduction technique
- The idea is that data may consists of n attributes. In the case of an image, this may mean n pixels
- PCA uses complicated maths to create an orthonormal basis with n attributes.

SUMMARY

- PCA essentially converts a data item of n attributes and calculates n different attributes to use.
- As much information as possible is put into the first attribute. As much information as possible is put into the second attribute
- This allows us to reduce the dimensionality of data by dropping as many of the least important attributes as we wish
- We can balance the accuracy of the compression with the reduction in data size as we choose
- There are many other useful unsupervised techniques

QUESTIONS



FURTHER READING

- K-means clustering: <https://www.analyticsvidhya.com/blog/2019/08/comprehensive-guide-k-means-clustering/>
- Autoencoder: <https://www.jeremyjordan.me/autoencoders/>
<https://www.v7labs.com/blog/autoencoders-guide>
- Principal Component Analysis (PCA): <https://builtin.com/data-science/step-step-explanation-principal-component-analysis>